



6CS030 Lecture 7

SQL on Hadoop

Hadoop Stack 'Zoo':

- Apache Spark

- Apache HBase

- Apache Pig

- Apache Hive

SQL on Hadoop

- MapReduce is very complex when compared to SQL
- Need for a more database-like setup on top of Hadoop
- Various projects can be used on top of Hadoop
 - See <http://hadoop.apache.org/> for a list.
- Sometimes referred to as a “Zoo”
- **ZooKeeper** provides a high-performance coordination service for distributed applications

SQL on Hadoop



■ HBase

□ <https://hbase.apache.org>



■ Pig

□ <https://pig.apache.org/>



■ Hive

□ <https://hive.apache.org/>



■ Spark

□ <https://spark.apache.org/>



HBase

- First Hadoop database inspired by Google's Bigtable
- Runs on top of HDFS
- NoSQL alike data storage platform
 - No typed columns, triggers, advanced query capabilities, etc.
- Offers a simplified structure and query language in a way that is highly scalable and can tackle large volumes



HBase

- Similar to RDBMSs, HBase organizes data in tables with rows and columns
- HBase table consists of multiple rows
- A row consists of a row key and one or more columns with values associated with them
- Rows in a table are sorted alphabetically by the row key

HBase

- Each column in HBase is denoted by a column family and qualifier (separated by a colon, ':')
- A column family physically co-locates a set of columns and their values
- Every row has the same column families, but not all column families need to have a value per row
- Each cell in a table is hence defined by a combination of the row key, column family and column qualifier, and a timestamp

HBase

- Example: HBase table to store and query users
- The row key will be the user id
- column families:qualifiers
 - name:first
 - name:last
 - email (without a qualifier)

HBase

```
hbase(main):001:0> create 'users', 'name', 'email'
0 row(s) in 2.8350 seconds
```

```
=> Hbase::Table - users
```

```
hbase(main):002:0> describe 'users'
Table users is ENABLED
users
COLUMN FAMILIES DESCRIPTION
{NAME => 'email', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
{NAME => 'name', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
2 row(s) in 0.3250 seconds
```

Examples from: www.pdbmbook.com

HBase

```
hbase(main):006:0> put 'users', 'seppe', 'name:first', 'Seppe'
```

```
0 row(s) in 0.0200 seconds
```

```
hbase(main):007:0> put 'users', 'seppe', 'name:last', 'vanden Broucke'
```

```
0 row(s) in 0.0330 seconds
```

```
hbase(main):008:0> put 'users', 'seppe', 'email', 'seppe.vandenbroucke@kuleuven'
```

```
0 row(s) in 0.0570 seconds
```

```
hbase(main):009:0> scan 'users'
```

ROW	COLUMN+CELL
seppe	column=email:, timestamp=1495293082872, value=seppe.vanden broucke@kuleuven.be
seppe	column=name:first, timestamp=1495293050816, value=Seppe
seppe	column=name:firstt, timestamp=1495293047100, value=Seppe
seppe	column=name:last, timestamp=1495293067245, value=vanden Broucke

```
1 row(s) in 0.1170 seconds
```

```
hbase(main):011:0> get 'users', 'seppe'
```

COLUMN	CELL
email:	timestamp=1495293082872, value=seppe.vandenbroucke@kuleuven.be
name:first	timestamp=1495293050816, value=Seppe
name:last	timestamp=1495293067245, value=vanden Broucke

```
4 row(s) in 0.1250 seconds
```

HBase

- HBase's query facilities are very limited
- Essentially a key-value, distributed data store with simple get/put operations
- Includes facilities to write MapReduce programs
- HBase (similar to Hadoop) doesn't perform well on less than 5 HDFS DataNodes with an additional NameNode
 - only makes the effort worthwhile when you can invest in, set up and maintain at least 6-10 nodes

Pig

- Yahoo! Developed “Pig”, which was made open source as Apache Pig in 2007
- High-level platform for creating programs that run on Hadoop which uses MapReduce underneath
 - The language used is Pig Latin
- Resembles the querying facilities of SQL

Pig

```
timesheet = LOAD 'timesheet.csv' USING PigStorage(',');
```

```
raw_timesheet = FILTER timesheet by $0 > 100;
```

```
timesheet_logged = FOREACH raw_timesheet GENERATE $0 AS  
driverId, $2 AS hours_logged, $3 AS miles_logged;
```

```
grp_logged = GROUP timesheet_logged by driverId;
```

```
sum_logged = FOREACH grp_logged GENERATE group as driverId,  
SUM(timesheet_logged.hours_logged) as sum_hourslogged,  
SUM(timesheet_logged.miles_logged) as sum_mileslogged;
```

Pig

- Some have argued that RDBMSs and SQL are substantially faster than MapReduce – and hence Pig
- Pig Latin is relatively procedural versus declarative SQL
- No wide adoption



Hive

- Initially developed by Facebook but open-sourced afterwards
- Data warehouse solution offering SQL querying facilities on top of Hadoop
- Converts SQL-like queries to a MapReduce pipeline
- Also offers a JDBC and ODBC interface
- Can run on top of HDFS, as well as other file systems

Hive

- Hive Metastore stores metadata for each table such as its schema and location on HDFS (using an RDBMS)
- Driver service is responsible to receive and handle incoming queries
 - query is first converted to an abstract syntax tree, which is then converted to a directed acyclic graph representing an execution plan
 - the directed acyclic graph will contain a number of MapReduce stages and tasks
- Optimizer optimizes the directed acyclic graph
- Executer sends MapReduce stages to Hadoop's resource manager (e.g. YARN) and monitor their progress

Hive

- HiveQL does not completely follow the full SQL-92 standard
 - E.g., lacks strong support for indexes, transactions, materialized views, and only has limited subquery support
- Example:
`SELECT genre, SUM(nrPages) FROM books GROUP BY genre`
- HiveQL also allows to query data sets other than structured tables

Hive

```
CREATE TABLE docs (line STRING); -- create a docs table
```

```
-- load in file from HDFS to docs table, overwrite existing data:
```

```
LOAD DATA INPATH '/testfile.txt' OVERWRITE INTO TABLE docs;
```

```
-- perform word count
```

```
SELECT word, count(1) AS count
```

```
FROM ( -- split each line in docs into words
```

```
    SELECT explode(split(line, '\s')) AS word FROM docs
```

```
) t
```

```
GROUP BY t.word
```

```
ORDER BY t.word;
```

Hive


- One difference with traditional RDBMS is that Hive does not enforce the schema at the time of loading the data
 - Hive: schema-on-read
 - RDBMS: schema-on-write
- Recent versions of Hive support full ACID transaction management
- Performance and speed of SQL queries still forms the main disadvantage of Hive today
 - Solutions to bypass MapReduce (e.g. Apache Tez, Cloudera Impala, Facebook Presto)

Apache Spark

- Open-source alternative for MapReduce
- New programming paradigm centered on a data structure called the **Resilient Distributed Dataset (RDD)**
 - RDDs can enable the construction of iterative programs that have to visit a data set multiple times, as well as more interactive or exploratory programs
 - RDD is a fundamental data structure of Spark
 - Each dataset in RDD is divided into logical partitions that can be computed on different nodes of the cluster
 - Is maintained in a fault tolerant way
 - RDDs can contain any type of Python, Java or Scala objects, including user-defined classes.
- Many orders of magnitude faster than MapReduce implementations
- Rapidly adopted by many Big Data vendors

Apache Spark

- Similar to Hadoop, Spark works with HDFS and requires a cluster manager (e.g. YARN)
- Key components
 - Spark Core
 - Spark SQL
 - MLib, Spark Streaming, GraphX



Spark Core

- Foundation for all other components
- Provides functionality for task scheduling and a set of basic data transformations
- Can be used through many programming languages
 - For example: Java, Python, Scala and R
- RDDs are the primary data abstraction in Spark
 - designed to support in-memory data storage and operations, distributed across a cluster
- Can be used to handle JSON and CSV data

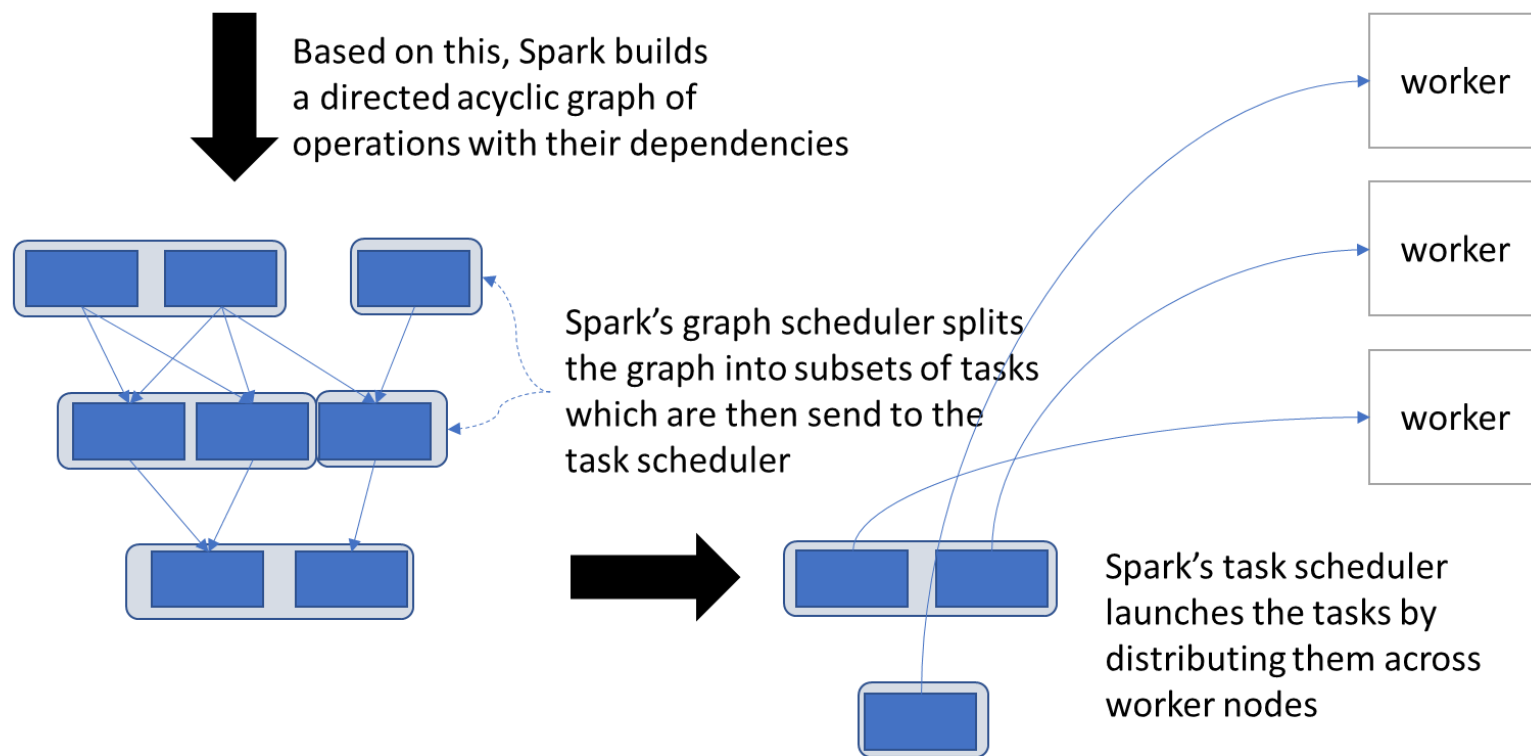
Spark Core

- Once data is loaded into an RDD, two basic types of operations can be performed:
 - Transformation which creates a new RDD through changing the original one
 - Actions which measure but do not change the original data
- Transformations are lazily evaluated
 - executed when a subsequent action has a need for the result
 - So can mean errors do not immediately appear
 - E.g., file does not exist
- RDDs will also be kept as long as possible in memory
- A chain of RDD operations gets compiled by Spark into a directed acyclic graph but which is then spread out and calculated over the cluster

Spark Core

A programmer writes a Spark program using its API:

```
rdc1.join(rdd2).groupBy(...).filter(...)
```



Spark Core – Map Reduce

- Spark's RDD API is relatively easy to work with compared to writing MapReduce programs

```
# Load in an RDD from a text file, the RDD will represent a  
# collection of text strings
```

```
 #(one for each line)
```

```
text_file = sc.textFile("testfile.txt")
```

Let's combine our
testfile1 and testfile2
cat test* > testfile.txt

```
# Count the word occurrences
```

```
counts = text_file.flatMap(lambda line: line.split("  
    .map(lambda word: (word, 1)) \  
    .reduceByKey(lambda a, b: a + b)
```

```
# counts is a PythonRDD
```

```
# Need loop to print items:
```

```
for x in counts.collect():  
    print x
```

```
(u'A', 1)  
(u'ago', 1)  
(u'episode', 1)  
(u'far', 2)  
(u'away', 1)  
(u'long', 1)  
(u'a', 1)  
(u'Another', 1)  
(u'Star', 1)  
(u'galaxy', 1)  
(u'of', 1)  
(u'in', 1)  
(u'Wars', 1)  
(u'time', 1)
```


Spark SQL

- Spark SQL runs on top of Spark Core and introduces another data abstraction called **DataFrames**
- DataFrames can be created from RDDs by specifying a schema on how to structure the data elements in the RDD, or can be loaded in directly from various sorts of file formats
- Even although DataFrames continue to use RDDs behind the scenes, they represent themselves to the end user as a collection of data organized into named columns
- You can run Spark directly using:
 - **pyspark** – uses Python
 - **spark-shell** – uses Scala
 - **spark-sql** – to run SQL queries
 - **spark-submit** – to run a program file, such as Python
- Or can access it via a programming language, such as Java
- See here for information:
 - <https://spark.apache.org/docs/latest/sql-programming-guide.html>

Spark SQL – Handling JSON (pyspark)

Create a DataFrame object by reading in a file

```
df = spark.read.json("student.json")
```

```
df.show()
```

```
+-----+-----+-----+-----+
| age|          course|          email|  name|
+-----+-----+-----+-----+
| null|BSc Horticulture|          null|   Tom|
|  45| MSc Agriculture|          null| Helen|
|  30|          null|S.Carter.borchest...| Alice|
|  21|BSc Horticulture|          null|Johnny|
+-----+-----+-----+-----+
```

student.json:

```
{ "name": "Tom",
  "course": "BSc Horticulture" }
{ "name": "Helen", "age": 45,
  "course": "MSc Agriculture" }
{ "name": "Alice", "age": 30,
  "email": "S.Carter@borchester.ac.uk" }
{ "name": "Johnny", "age": 21,
  "course": "BSc Horticulture" }
```

DataFrames are structured in columns and rows:

```
df.printSchema()
```

```
root
 |-- age: long (nullable = true)
 |-- course: string (nullable = true)
 |-- email: string (nullable = true)
 |-- name: string (nullable = true)
```

Spark SQL (pyspark)

```
df.select("name").show()
```

```
+-----+  
|  name |  
+-----+  
|   Tom |  
|  Helen |  
|  Alice |  
|Johnny |  
+-----+
```

SQL-like operations can now easily be expressed:

```
df.select(df['name'], df['age'] + 1).show()
```

```
+-----+-----+  
|  name|(age + 1)|  
+-----+-----+  
|   Tom|        null|  
|  Helen|         46|  
|  Alice|         31|  
|Johnny|         22|  
+-----+-----+
```

Spark SQL (pyspark)

```
df.filter(df['age'] > 21).show()
```

```
--+-----+-----+-----+
|age|      course|      email| name|
+---+-----+-----+-----+
| 45|MSc Agriculture|      null|Helen|
| 30|      null|S.Carter.borchest...|Alice|
+---+-----+-----+-----+
```

```
df.groupBy("course").count().show()
```

```
+-----+-----+
|      course|count|
+-----+-----+
|MSc Agriculture|    1|
|      null|    1|
|BSc Horticulture|    2|
+-----+-----+
```

Spark SQL (pyspark)

- Spark implements a full SQL query engine which can convert SQL statements to a series of RDD transformations and actions
- First Register the DataFrame as a SQL temporary view:
`df.createOrReplaceTempView("student")`

Can then use SQL like syntax:

```
sqlDF = spark.sql("SELECT name, age, course FROM student WHERE  
age > 21")
```

```
sqlDF.show()
```

```
+-----+-----+  
| name|age|          course|  
+-----+-----+  
|Helen| 45|MSc Agriculture|  
|Alice| 30|          null|  
+-----+-----+
```

You can not just type
the SQL Code at the
command line

See Workbook for
further examples

MLlib, Spark Streaming and GraphX

- There are lots of other Spark tools to help work with Big Data:
 - **MLlib** is Spark's machine learning library
 - offers classification, regression, clustering, and recommender system algorithms
 - **Spark Streaming** uses Spark Core and its scheduling engine to perform streaming analytics
 - provides a high-level concept called **DStream**, which represents a continuous stream of data
 - **GraphX** is Spark's component implementing programming abstractions to deal with graph based structures
 - based on the RDD abstraction
 - comes with a set of fundamental operators and algorithms to work with graphs and simplify graph analytics tasks

MLlib, Spark Streaming and GraphX

■ Example: Word counting

```
from pyspark import SparkContext
from pyspark.streaming import StreamingContext
sc = SparkContext("local[2]", "StreamingWordCount")
ssc = StreamingContext(sc, 1)

# Create a DStream that will connect to server.mycorp.com:9999 as a source
lines = ssc.socketTextStream("server.mycorp.com ", 9999)

# Split each line into words
words = lines.flatMap(lambda line: line.split(" "))

# Count each word in each batch
pairs = words.map(lambda word: (word, 1))
wordCounts = pairs.reduceByKey(lambda x, y: x + y)

# Print out first ten elements of each RDD generated in the wordCounts Dstream
wordCounts.pprint()

# Start the computation
ssc.start()
ssc.awaitTermination()
```

Conclusion

- This lecture has looked at:
 - SQL on Hadoop
 - Apache Spark
- This week's workbook will look at using:
 - Further Hadoop examples
 - How to use CSV files
 - Apache Spark
 - SQL queries
 - JSON and CSV data